AUDIT REPORT

intermittent-demand-forecasting: demo video audit report

A slide-by-slide record of every number, assumption and calculation in the demo video, with its source.

Item	Value
Repository	https://github.com/kavya2693/intermittent-demand-forecasting
Numbers taken from commit	`3616b6c2747b6be20b0c56b6345c4954391b22ac` (outputs last changed in `3616b6c (2026-10-13 22:03:11 +0400)`)
Video	`docs/demo/demo.mp4` · 88.4 s · 1920x1080 · 30 fps · silent
Generated	2026-10-14 14:20
Automated checks	63 pass · 0 fail · 1 manual (of 64 figures)
Assumptions	12 (A-01 to A-12)

How to use this report

Every number or claim on screen has an ID (F-01, F-02, ...) and every modelling assumption has an ID (A-01, ...). Cite the ID. Each figure lists its source file, the exact column, row or key, the formula for derived numbers, and an automated check result. The change-impact matrix (section 5) shows which slides and files change if an input or assumption changes.

Types: direct = read straight from a repo output; derived = computed from direct values (formula given); illustrative = drawing or example, not data; general knowledge = context not produced by the repo.

Regenerate

```
make run                                    # outputs/*.csv, run_report.md, forecasts.parquet (~12 min)
PYTHONPATH=src .venv/bin/python docs/demo/source/tools/export_forecast_sheet.py
docs/demo/source/assets/forecast_sheet.json
python3 docs/demo/source/tools/inject_sheet.py docs/demo/source/assets/forecast_sheet.json
docs/demo/source/compositions/value.html --block SHEET --keys
origin_week,window_end_week,skus,model,wape_8wk_all,forecast_units_8wk,sold_units_8wk,rows
python3 docs/demo/source/tools/inject_sheet.py docs/demo/source/assets/forecast_sheet.json
docs/demo/source/compositions/s04-routes.html --block RUN --keys
scatter,example,forecast_rows_total,forecast_models
python3 docs/demo/source/tools/make_end_scatter.py outputs/sku_segments.csv docs/demo/source/assets/scatter-
end.png
cd docs/demo/source && npx --yes hyperframes@0.8.156 check && npx --yes hyperframes@0.8.156 render -f 30 -q
looks -o renders/intermittent-demand-forecasting-master.mp4 && cd -
ffmpeg -y -i docs/demo/source/renders/intermittent-demand-forecasting-master.mp4 -c:v libx264 -preset slow
-crf 25 -pix_fmt yuv420p -movflags +faststart -an docs/demo/demo.mp4
python3 docs/demo/source/tools/build_audit_report.py docs/demo/source/audit/facts.yaml   # stills, checks,
MD/DOCX/PDF
```

1. Figures

ID	Shown	Slides	Source	Column / row / key	How computed	Type	Check
F-01	3 industry chips	01	author's framing	-	-	general knowledge	manual: Context statement, not a repo output.
F-02	4 problem cards	01	README 'The problem'	README lines 14-28	-	general knowledge	✓ verified
F-03	1,500 SKUs	02, 05, 06, 07, 08	outputs/sku_segments.csv	row count	-	direct	✓ verified
F-04	156 weeks / 1,500 × 156 wk	02, 05	outputs/run_log.txt (printed by cli.py from data/sales.parquet)	'1500 SKUs, 156 weeks'	-	direct	✓ verified
F-05	670 are slow movers	02	outputs/sku_segments.csv	adi	count(adi > 2) = 670	derived	✓ verified
F-06	slow = sells in < half the weeks	02	src/demand_segments/segments.py:26 (adi = len(h) / len(nz))	ADI = observed weeks / weeks with a sale	ADI > 2 <=> sale weeks < half of observed weeks	direct	✓ verified
F-07	7 methods, same windows / Seven methods	02, 05	outputs/forecasts.parquet (model column)	models minus ZeroForecast and LightGBM-history-only	9 - 2 = 7	direct	✓ verified
F-08	weekly WAPE crowns zero: 829 SKUs	02, 04	outputs/segment_counts.csv; outputs/scores_by_segment.csv	skus where segment in {intermittent, lumpy}	397 + 432 = 829	derived	✓ verified
F-09	crowns zero / zero beats every model	02, 04	outputs/scores_by_segment.csv	argmin(wape) per segment	segments whose best weekly WAPE is ZeroForecast = {intermittent, lumpy}	derived	✓ verified
F-10	method per type · 8-week totals · losses shown	03	outputs/champions.csv; outputs/scores_by_segment.csv (wape_8wk); outputs/lightgbm_vs_best_statistical.csv	4 champion rows; wape_8wk column; win share < 50% on intermittent	-	direct	✓ verified
F-11	REPLAY THE PAST · step forward, repeat · refit ×6	03, 05	src/demand_segments/backtest.py:29-31; forecast_stats.py:49-57	origins(); cross_validation(step_size=HORIZON)	-	direct	✓ verified
F-12	SOLD 0 0 0 4 0 0 0 4 · ZERO · STEADY · winners	04	tests/test_metrics.py:70-80 (test_eight_week_wape_stops_rewarding_the_zero_forecast)	y = tile([0,0,0,4], 2)	weekly WAPE: steady 12/8 = 150% > zero 100%; 8-week WAPE: steady 0%, zero 100%	illustrative	✓ verified
F-14	sales.parquet, skus.parquet, calendar.py, data.py, segments.py, features.py, forecast_stats.py, forecast_ml.py, backtest.py, metrics.py, champions(), report.py, fold_summary.py, tests/ · make check	05	src/demand_segments/, scripts/, tests/, Makefile	files + defs	-	direct	✓ verified
F-16	8 categories	05	src/demand_segments/data.py:20-29 (CATEGORIES)	dict entries	-	direct	✓ verified
F-17	data.py · seed 7	05	src/demand_segments/data.py:39	seed: int = 7	-	direct	✓ verified
F-18	ADI 1.32 · CV² 0.49	05, 06	src/demand_segments/segments.py:14-15	ADI_CUT, CV2_CUT	-	direct	✓ verified
F-19	17 origin + 13 target	05	src/demand_segments/features.py (ORIGIN_FEATURES, TARGET_FEATURES) via demo assets/forecast_sheet.json arch	len(ORIGIN_FEATURES), len(TARGET_FEATURES)	-	direct	✓ verified
F-20	300k–539k rows / fit	05	features.training_frame(panel, origin) via demo assets/forecast_sheet.json arch	train_rows_first_origin = 299,648; train_rows_last_origin = 539,390	rounded to thousands	derived	✓ verified
F-21	SNaive52 · MA13 · ETS · Croston · TSB · ARIMA	05	src/demand_segments/forecast_stats.py:25-33	models()	-	direct	✓ verified
F-22	329 s · 4 procs	05	outputs/run_report.md; Makefile:8	'statistical models: 329'; --jobs 4	-	direct	✓ verified
F-23	Tweedie · 400 rounds · direct, h = 1..8 · 256 s · 4 threads	05	src/demand_segments/forecast_ml.py:14, :25, :27; features.py:21; outputs/run_report.md	objective, num_threads, ROUNDS, HORIZON, 'LightGBM: 256'	-	direct	✓ verified
F-24	Ablation · history only · 144 s	05	outputs/run_report.md	'LightGBM-history-only: 144'	-	direct	✓ verified
F-25	same 48 weeks for every model · origins, 8 weeks apart · 6 / 6	05, 06	outputs/forecasts.parquet	distinct weeks; 6 folds	6 × 8 = 48	direct	✓ verified

ID	Shown	Slides	Source	Column / row / key	How computed	Type	Check
F-26	✕ misaligned windows → stop	05	src/demand_segments/backtest.py:56-59	RuntimeError	-	direct	✓ verified
F-27	WAPE · 8-wk WAPE · bias · MAPE	05	src/demand_segments/metrics.py	wape, window_wape, bias, mape	-	direct	✓ verified
F-28	pick on origins 1–3, grade 4–6	05, 08	src/demand_segments/report.py:30; backtest.py:109	n_select = n_windows // 2 = 3	-	direct	✓ verified
F-29	✕ zero, ablation: not eligible	05	src/demand_segments/backtest.py:108; tests/test_backtest.py:41	isin([ZERO, ML_HISTORY_ONLY])	-	direct	✓ verified
F-30	8 CSV + run_report	05	src/demand_segments/report.py:41-48	.to_csv calls	-	direct	✓ verified
F-31	wins · |bias| > 10%	05	scripts/fold_summary.py	origins_with_abs_bias_over_10pct	-	direct	✓ verified
F-32	CI guards · leakage test · metric pins · same windows · no 0 champion	05	tests/*.py; .github/workflows/ci.yml; Makefile check	test names	-	direct	✓ verified
F-33	make run ≈ 12 min / ~12 min	05, 10	outputs/run_report.md 'Run time (seconds)'	329 + 256 + 144 = 729 s	729 / 60 = 12.2 min	derived	✓ verified
F-34	batch, offline · no API key or model server	05	README 'Running it'	'Nothing needs an API key or a model server'	-	direct	✓ verified
F-35	champion path: LightGBM won all 4 segments	05	outputs/champions.csv	champion	count(champion == LightGBM)	direct	✓ verified
F-36	refit ×6	05	src/demand_segments/backtest.py:68-71 (fit per origin); forecast_stats.py:49-57 (cross_validation n_windows)	for fold, origin in enumerate(starts)	-	direct	✓ verified
F-37	325 / 346 / 397 / 432	06	outputs/segment_counts.csv	skus	-	direct	✓ verified
F-38	1,500 dots on the ADI / CV² plane	06	outputs/sku_segments.csv; RUN block in compositions/s04-routes.html	adi, cv2, segment (rounded to 3 decimals)	-	direct	✓ verified
F-39	3.9/3 · 3.7/3 · 3.6/4 · 3.2/1 · 4.0/1 · 2.9/5	06	outputs/forecasts.parquet	sku_id = SKU00191, model = LightGBM, sum(yhat), sum(y) per fold	sums per fold, forecast rounded to 0.1	derived	✓ verified
F-40	about 0.4 a week	06	outputs/forecasts.parquet	SKU00191, LightGBM, mean(yhat) over 48 weeks = 0.44	-	derived	✓ verified
F-41	9 models × 1,500 SKUs × 48 weeks = 648,000 forecasts	06	outputs/forecasts.parquet	row count	9 × 1,500 × 48	direct	✓ verified
F-42	One spare part	06	outputs/sku_segments.csv; data/skus.parquet	SKU00191	-	direct	✓ verified
F-43	SKU01050, SKU01018, SKU00191, SKU00029, SKU01486, SKU00330	07	outputs/forecasts.parquet (fold 6, LightGBM); outputs/sku_segments.csv; data/skus.parquet; data/sales.parquet; SHEET block in compositions/value.html	per-SKU sums	recomputed from the repo files, compared to the injected SHEET	direct	✓ verified
F-44	+ 1,494 more SKUs on this sheet	07	F-03	1,500 - 6 rows	1,500 - 6 = 1,494	derived	✓ verified
F-45	215,538 forecast vs 222,039 sold	07	outputs/forecasts.parquet	fold 6, LightGBM: sum(yhat), sum(y)	-	derived	✓ verified
F-46	8-week WAPE 20.7%	07	outputs/forecasts.parquet	fold 6, LightGBM	sum |8-wk forecast - 8-wk sold| / sum 8-wk sold over SKUs (metrics.window_wape)	derived	✓ verified
F-47	FROM 2025-11-03 · LightGBM	07	outputs/forecasts.parquet	first week of fold 6	-	direct	✓ verified
F-50	21.9%	08	outputs/scores_by_segment.csv	segment=all, model=LightGBM, wape_8wk	x 100	direct	✓ verified
F-51	seasonal naive 29.1%	08	outputs/scores_by_segment.csv	segment=all, model=SeasonalNaive, wape_8wk	x 100	direct	✓ verified
F-52	-25.8%	08	outputs/champions.csv	segment=smooth, wape_reduction_vs_seasonal_naive	-(x 100)	direct	✓ verified
F-53	37.6% → 27.9%	08	outputs/champions.csv	seasonal_naive_wape, champion_wape	-	direct	✓ verified
F-54	-30.7%	08	outputs/champions.csv	segment=erratic, wape_reduction_vs_seasonal_naive	-(x 100)	direct	✓ verified
F-55	74.8% → 51.9%	08	outputs/champions.csv	seasonal_naive_wape, champion_wape	-	direct	✓ verified
F-56	24/24 · 4 segments ×	08	outputs/fold_summary.csv	sum(origins_won_on_w	-	direct	✓ verified

ID	Shown	Slides	Source	Column / row / key	How computed	Type	Check
	6 origins			ape_8wk), sum(origins)			
F-57	LightGBM wins all four segments on the 8-week total.	08	outputs/scores_by_segment.csv	argmin(wape_8wk) excluding ZeroForecast and the ablation	-	derived	✓ verified
F-58	100.0%	09	outputs/scores_by_segment.csv	segment=intermittent, model=ZeroForecast, wape	x 100	direct	✓ verified
F-59	121.1%	09	outputs/scores_by_segment.csv	segment=intermittent, model=LightGBM, wape	x 100	direct	✓ verified
F-60	121.9%	09	outputs/scores_by_segment.csv; outputs/lightgbm_vs_best_statistical.csv	segment=intermittent, model=AutoARIMA, wape; best_statistical = AutoARIMA	x 100	direct	✓ verified
F-61	122.9%	09	outputs/scores_by_segment.csv	segment=intermittent, model=LightGBM-history-only, wape	x 100	direct	✓ verified
F-62	397 SKUs	09	outputs/segment_counts.csv	segment=intermittent	-	direct	✓ verified
F-63	45.6%	09	outputs/lightgbm_vs_best_statistical.csv	segment=intermittent, lightgbm_sku_win_share	x 100	direct	✓ verified
F-64	2 of 6 origins	09	outputs/fold_summary.csv	segment=intermittent, origins_won_on_wape / origins	-	direct	✓ verified
F-65	the edge is the calendar and promo plan, not the algorithm	09	outputs/scores_by_segment.csv	intermittent: LightGBM-history-only 122.9% > AutoARIMA 121.9% > LightGBM 121.1%	-	derived	✓ verified
F-66	MIT	10	LICENSE line 1	-	-	direct	✓ verified
F-67	$ make run # 1,500 SKUs, 6 origins, ~12 min	10	Makefile:7-9; README 'Running it'	run target	-	direct	✓ verified
F-68	LightGBM · statsforecast · pandas · Python	10	pyproject.toml dependencies	lightgbm, statsforecast, pandas	-	direct	✓ verified

2. Automated check run (2026-10-14 14:20)

ID	Expected on slide	Recomputed from repo	Result
F-01	3 industry chips	-	manual
F-02	true	true	✓ pass
F-03	1,500	1,500	✓ pass
F-04	true	true	✓ pass
F-05	670	670	✓ pass
F-06	true	true	✓ pass
F-07	7	7	✓ pass
F-08	829	829	✓ pass
F-09	true	true	✓ pass
F-10	true	true	✓ pass
F-11	true	true	✓ pass
F-12	true	true	✓ pass
F-14	true	true	✓ pass
F-16	8	8	✓ pass
F-17	true	true	✓ pass
F-18	true	true	✓ pass
F-19	17 origin + 13 target	17 origin + 13 target	✓ pass
F-20	300k–539k	300k–539k	✓ pass
F-21	true	true	✓ pass
F-22	true	true	✓ pass
F-23	true	true	✓ pass
F-24	true	true	✓ pass
F-25	true	true	✓ pass
F-26	true	true	✓ pass
F-27	true	true	✓ pass
F-28	true	true	✓ pass
F-29	true	true	✓ pass
F-30	8	8	✓ pass
F-31	true	true	✓ pass
F-32	true	true	✓ pass
F-33	~12 min	~12 min	✓ pass
F-34	true	true	✓ pass
F-35	4	4	✓ pass
F-36	true	true	✓ pass
F-37	325/346/397/432	325/346/397/432	✓ pass
F-38	true	true	✓ pass
F-39	3.9/3 · 3.7/3 · 3.6/4 · 3.2/1 · 4.0/1 · 2.9/5	3.9/3 · 3.7/3 · 3.6/4 · 3.2/1 · 4.0/1 · 2.9/5	✓ pass
F-40	0.4	0.4	✓ pass

ID	Expected on slide	Recomputed from repo	Result
F-41	648,000	648,000	✓ pass
F-42	true	true	✓ pass
F-43	true	true	✓ pass
F-44	1,494	1,494	✓ pass
F-45	215,538 forecast vs 222,039 sold	215,538 forecast vs 222,039 sold	✓ pass
F-46	20.7%	20.7%	✓ pass
F-47	2025-11-03	2025-11-03	✓ pass
F-50	21.9%	21.9%	✓ pass
F-51	29.1%	29.1%	✓ pass
F-52	-25.8%	-25.8%	✓ pass
F-53	37.6% → 27.9%	37.6% → 27.9%	✓ pass
F-54	-30.7%	-30.7%	✓ pass
F-55	74.8% → 51.9%	74.8% → 51.9%	✓ pass
F-56	24/24	24/24	✓ pass
F-57	true	true	✓ pass
F-58	100.0%	100.0%	✓ pass
F-59	121.1%	121.1%	✓ pass
F-60	true	true	✓ pass
F-61	122.9%	122.9%	✓ pass
F-62	397	397	✓ pass
F-63	45.6%	45.6%	✓ pass
F-64	2 of 6	2 of 6	✓ pass
F-65	true	true	✓ pass
F-66	true	true	✓ pass
F-67	true	true	✓ pass
F-68	true	true	✓ pass

3. Assumptions register

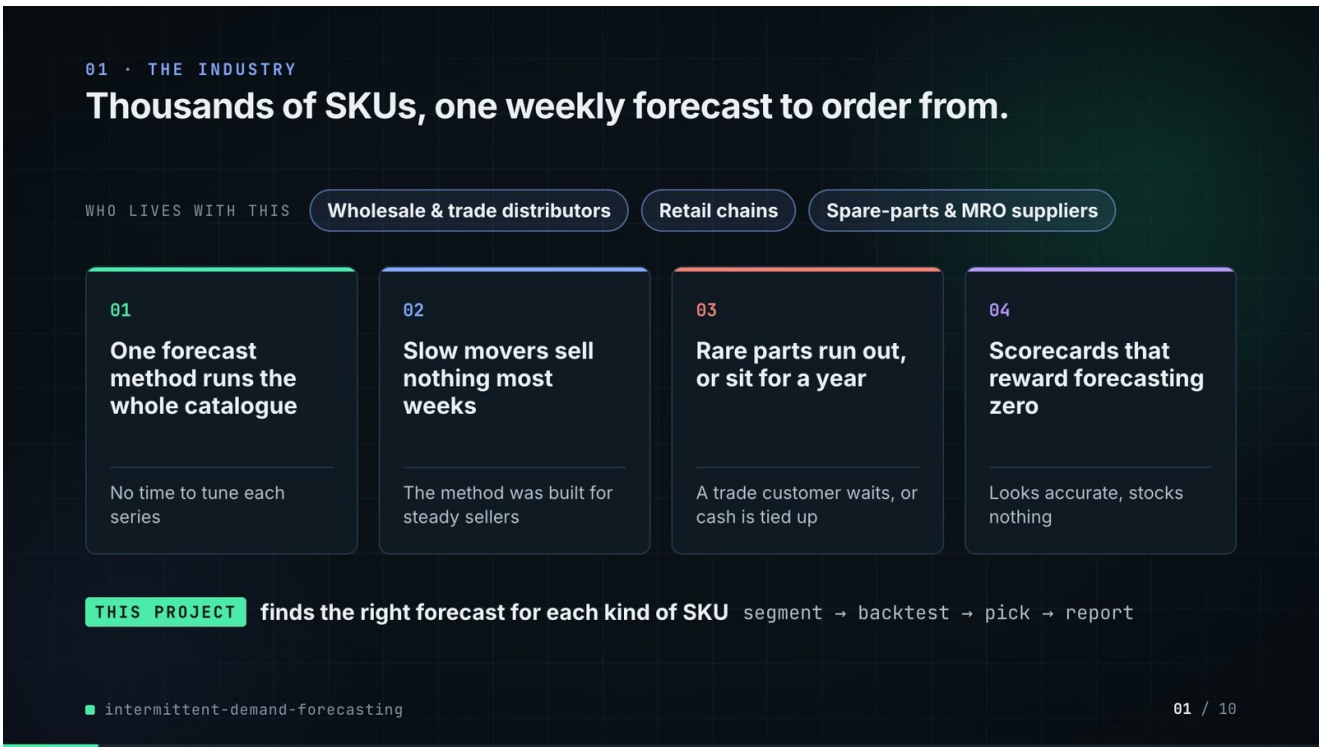
ID	Assumption	Source	Figures affected	Slides affected
A-01	All data is synthetic, produced by a seeded generator sized to the production problem: 1,500 SKUs, 156 weekly periods, 8 categories, B2B and B2C channels, seed 7. The production data it stands in for is proprietary; the patterns are the ones written into the generator.	src/demand_segments/data.py:20-29 (CATEGORIES), :37-39 (GeneratorConfig n_skus=1500, n_weeks=156, seed=7); README 'The data'	all figures	slides 01, 02, 03, 04, 05, 06, 07, 08, 09, 10
A-02	Segments use the Syntetos-Boylan cut-offs ADI 1.32 and CV² 0.49, computed on the history before the first origin, and stay fixed for all six origins.	src/demand_segments/segments.py:14-15; src/demand_segments/backtest.py:51	F-05, F-06, F-08, F-09, F-18, F-37, F-38, F-62	slides 02, 04, 05, 06, 09
A-03	Walk-forward backtest: six origins, eight weeks apart, each forecasting the next eight weeks (48 test weeks), the same windows for every model.	src/demand_segments/cli.py:18 (--windows default 6); features.py:21 (HORIZON = 8); backtest.py:29-31 (origins)	F-07, F-11, F-25, F-36, F-39, F-41, F-56, F-64	slides 02, 03, 05, 06, 08, 09
A-04	Each segment's champion is picked on eight-week WAPE over origins 1-3 and graded on origins 4-6; the zero forecast and the history-only ablation are never eligible.	src/demand_segments/report.py:30 (n_select = n_windows // 2); backtest.py:100-130 (champions), :108	F-10, F-28, F-29, F-35, F-52, F-53, F-54, F-55	slides 03, 05, 08
A-05	ETS and ARIMA are fitted without a 52-week season (season_length=1); only seasonal naive and LightGBM carry the yearly pattern, so the statistical side is weaker on seasonal SKUs than a tuned seasonal model would be.	src/demand_segments/forecast_stats.py:22-33; README 'Known limitations'	F-21, F-53, F-55, F-58, F-59, F-60, F-61, F-63, F-64	slides 05, 08, 09
A-06	LightGBM settings: Tweedie objective (variance power 1.3), 400 rounds, 63 leaves, seed 7, 4 threads; one global model, direct 1-8 week horizon, target = sales / SKU level.	src/demand_segments/forecast_ml.py:11-27	F-23, F-50, F-52, F-54, F-56, F-57, F-59	slides 05, 08, 09
A-07	LightGBM may read the calendar, the promotion plan and the price list for the target week (11 'known-future' features), because a planner has them in advance. The per-SKU statistical models see sales history only.	src/demand_segments/features.py:1-9 (docstring), forecast_ml.py:30-42 (KNOWN_FUTURE)	F-24, F-61, F-65	slides 05, 09
A-08	Feed-gap weeks are filled with a trailing 8-week mean for the statistical models' inputs only; every score uses the raw actuals, and missing actual weeks are skipped, never scored as zero.	src/demand_segments/forecast_stats.py:37-46 (impute); metrics.py:9-14 (_clean)	F-41, F-45, F-46	slides 06, 07
A-09	Run times are wall-clock seconds from one make run on the author's laptop with --jobs 4 (statistical models) and 4 LightGBM threads; they change with the machine and its load.	outputs/run_report.md 'Run time (seconds)'; Makefile:8 (--jobs 4); README 'What I got wrong first'	F-22, F-23, F-24, F-33, F-67	slides 05, 10
A-10	Slide 07 shows the LAST backtest origin (2025-11-03, weeks to 2025-12-22) as if it were the planner's sheet, with what actually sold next to it. Rows are the SKU at the median 8-week absolute % error of each segment (plus the next one in the two slow segments), not the best cases. The stock-room drain	docs/demo/source/tools/export_forecast_sheet.py (row selection); README 'Known limitations'	F-43, F-44, F-45, F-46, F-47	slides 07

ID	Assumption	Source	Figures affected	Slides affected
	is forecast total minus weekly actual sales, an illustration, not an inventory simulation (README: point forecasts only).			
A-11	Video rounding: percentages to 1 decimal, units and counts to whole numbers with thousands separators, 8-week totals under 100 units to 1 decimal; '300k-539k', '≈ 12 min' and 'about 0.4 a week' are rounded labels.	this audit (display convention)	F-20, F-33, F-40, F-45, F-46, F-50, F-51, F-52, F-53, F-54, F-55, F-58, F-59, F-60, F-61, F-63, F-67	slides 05, 06, 07, 08, 09, 10
A-12	Slide 01 (industries, everyday problems) is general demand-planning knowledge used as context. It is not a result of this repo and carries no numbers.	author's framing; README 'The problem' supports problems 1, 2 and 4	F-01, F-02	slides 01

4. Slide by slide

Slide 01: The industry

0.0-7.2 s · composition `compositions/s01-context.html`



On-screen text (verbatim)

01 · THE INDUSTRY
Thousands of SKUs, one weekly forecast to order from.
WHO LIVES WITH THIS · Wholesale & trade distributors · Retail chains · Spare-parts & MRO suppliers
01 One forecast method runs the whole catalogue — No time to tune each series
02 Slow movers sell nothing most weeks — The method was built for steady sellers
03 Rare parts run out, or sit for a year — A trade customer waits, or cash is tied up
04 Scorecards that reward forecasting zero — Looks accurate, stocks nothing
THIS PROJECT finds the right forecast for each kind of SKU · segment → backtest → pick → report

Figures on this slide

ID	Shown	Source · key	Calculation	Check
F-01	3 industry chips	`author's framing` · -	-	manual
F-02	4 problem cards	`README 'The problem'` · README lines 14-28	-	✓

Assumptions: A-12

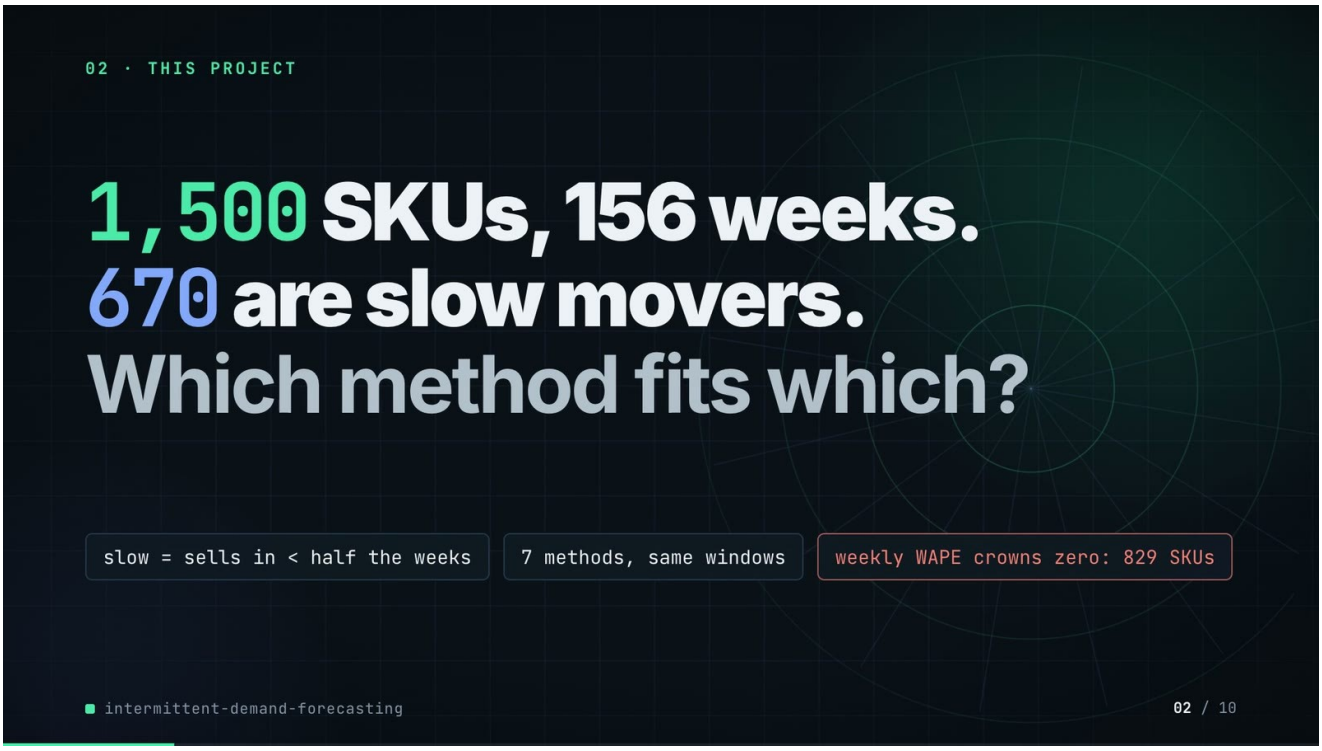
Illustrative, not data:

- The whole slide is context; no repo figures.

To change this slide: Edit the chips (#ctx-c1..3) and cards (#ctx-k1..4) in compositions/s01-context.html.

Slide 02: This project

6.8-12.0 s · composition `compositions/s02-hook.html`



On-screen text (verbatim)

02 · THIS PROJECT

1,500 SKUs, 156 weeks. 670 are slow movers. Which method fits which?

slow = sells in < half the weeks · 7 methods, same windows · weekly WAPE crowns zero: 829 SKUs

Figures on this slide

ID	Shown	Source · key	Calculation	Check
F-03	1,500 SKUs	`outputs/sku_segments.csv` · row count	-	✓
F-04	156 weeks / 1,500 × 156 wk	`outputs/run_log.txt (printed by cli.py from data/sales.parquet)` · '1500 SKUs, 156 weeks'	-	✓
F-05	670 are slow movers	`outputs/sku_segments.csv` · adi	count(adi > 2) = 670	✓
F-06	slow = sells in < half the weeks	`src/demand_segments/segments.py:26 (adi = len(h) / len(nz))` · ADI = observed weeks / weeks with a sale	ADI > 2 <=> sale weeks < half of observed weeks	✓
F-07	7 methods, same windows / Seven methods	`outputs/forecasts.parquet (model column)` · models minus ZeroForecast and LightGBM-history-only	9 - 2 = 7	✓
F-08	weekly WAPE crowns zero: 829 SKUs	`outputs/segment_counts.csv; outputs/scores_by_segment.csv` · skus where segment in {intermittent, lumpy}	397 + 432 = 829	✓
F-09	crowns zero / zero beats every model	`outputs/scores_by_segment.csv` · argmin(wape) per segment	segments whose best weekly WAPE is ZeroForecast = {intermittent, lumpy}	✓

Assumptions: A-01, A-02, A-03

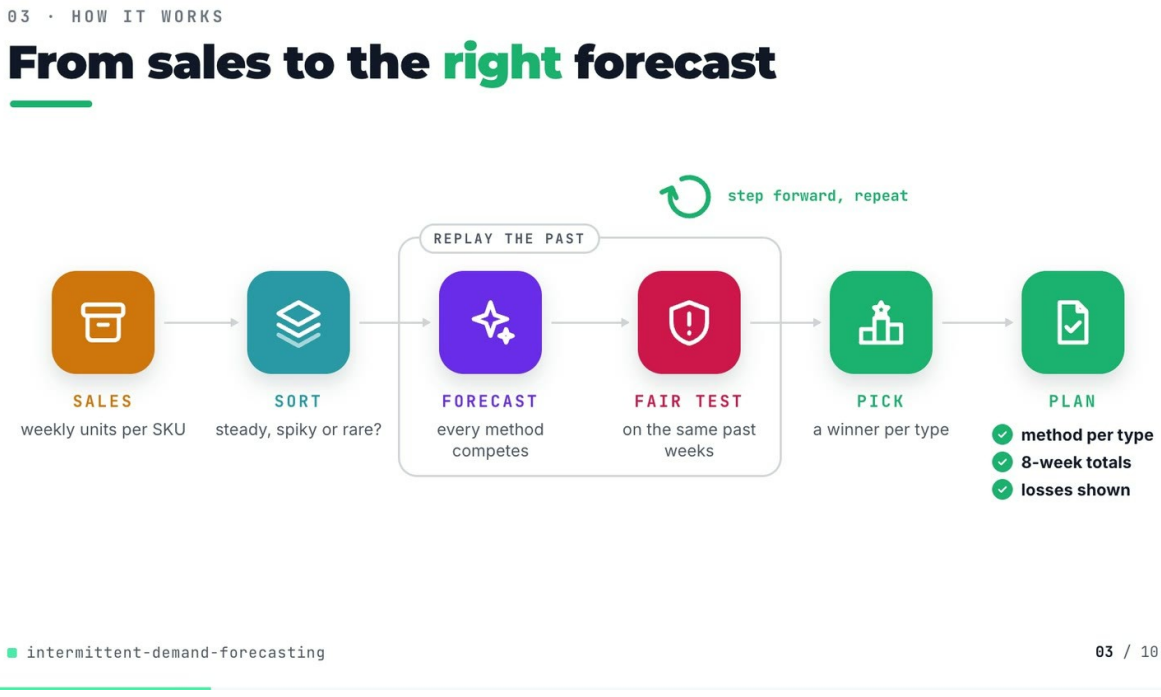
Illustrative, not data:

- Rings in the background are decoration.

To change this slide: Headline spans #hook-l1..l3 and chips #hook-f1..f3 in compositions/s02-hook.html.

Slide 03: How it works (white chapter, slide A)

11.7-19.7 s · composition `compositions/explainer.html (CONFIG.blocks, CONFIG.groups, CONFIG.loop)`



On-screen text (verbatim)

03 · HOW IT WORKS

From sales to the right forecast

SALES weekly units per SKU → SORT steady, spiky or rare? → [REPLAY THE PAST: FORECAST every method competes → FAIR TEST on the same past weeks] → PICK a winner per type → PLAN ✓ method per type ✓ 8-week totals ✓ losses shown

step forward, repeat (loop over the fair test)

Figures on this slide

ID	Shown	Source · key	Calculation	Check
F-10	method per type · 8-week totals · losses shown	`outputs/champions.csv`; `outputs/scores_by_segment.csv (wape_8wk)`; `outputs/lightgbm_vs_best_statistical.csv` · 4 champion rows; wape_8wk column; win share < 50% on intermittent	-	✓
F-11	REPLAY THE PAST · step forward, repeat · refit ×6	`src/demand_segments/backtest.py:29-31`; `forecast_stats.py:49-57` · origins(); cross_validation(step_size=HORIZON)	-	✓

Assumptions: A-02, A-03, A-04

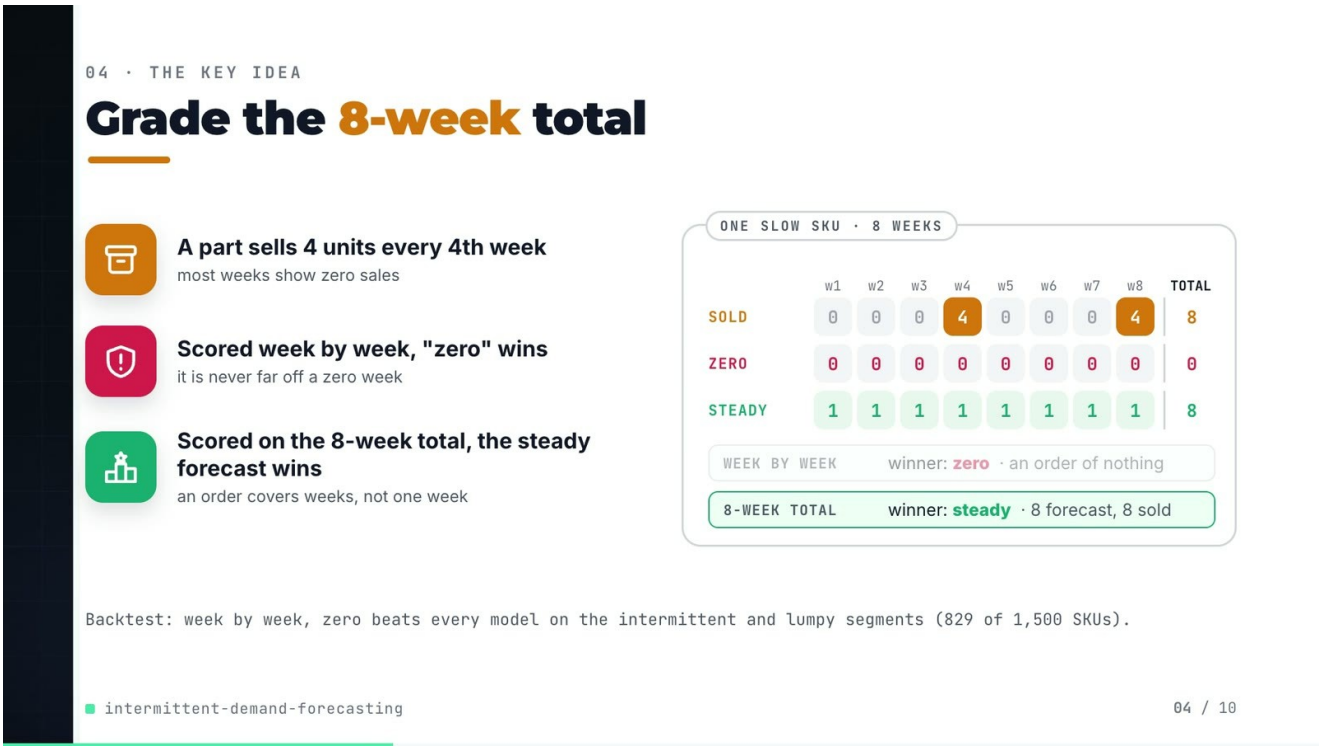
Illustrative, not data:

- Icons and colours are a visual language, not data.
- 'Type' is the plain-words name for a demand segment.

To change this slide: Edit CONFIG in compositions/explainer.html; badge truth is F-10.

Slide 04: The key idea (white chapter, slide B)

19.7-26.7 s · composition `compositions/explainer.html (#ex-b ...)`



On-screen text (verbatim)

04 · THE KEY IDEA

Grade the 8-week total

A part sells 4 units every 4th week — most weeks show zero sales

Scored week by week, "zero" wins — it is never far off a zero week

Scored on the 8-week total, the steady forecast wins — an order covers weeks, not one week

ONE SLOW SKU · 8 WEEKS: SOLD 0 0 0 4 0 0 0 4 (total 8) · ZERO 0 ×8 (0) · STEADY 1 ×8 (8)

WEEK BY WEEK winner: zero · an order of nothing | 8-WEEK TOTAL winner: steady · 8 forecast, 8 sold

Backtest: week by week, zero beats every model on the intermittent and lumpy segments (829 of 1,500 SKUs).

Figures on this slide

ID	Shown	Source · key	Calculation	Check
F-08	weekly WAPE crowns zero: 829 SKUs	`outputs/segment_counts.csv`; `outputs/scores_by_segment.csv` · skus where segment in {intermittent, lumpy}	397 + 432 = 829	✓
F-09	crowns zero / zero beats every model	`outputs/scores_by_segment.csv` · argmin(wape) per segment	segments whose best weekly WAPE is ZeroForecast = {intermittent, lumpy}	✓
F-12	SOLD 0 0 0 4 0 0 0 4 · ZERO · STEADY · winners	`tests/test_metrics.py:70-80 (test_eight_week_wape_stops_rewarding_the_zero_forecast)` · y = tile([0,0,0,4], 2)	weekly WAPE: steady 12/8 = 150% > zero 100%; 8-week WAPE: steady 0%, zero 100%	✓

Assumptions: A-02, A-04

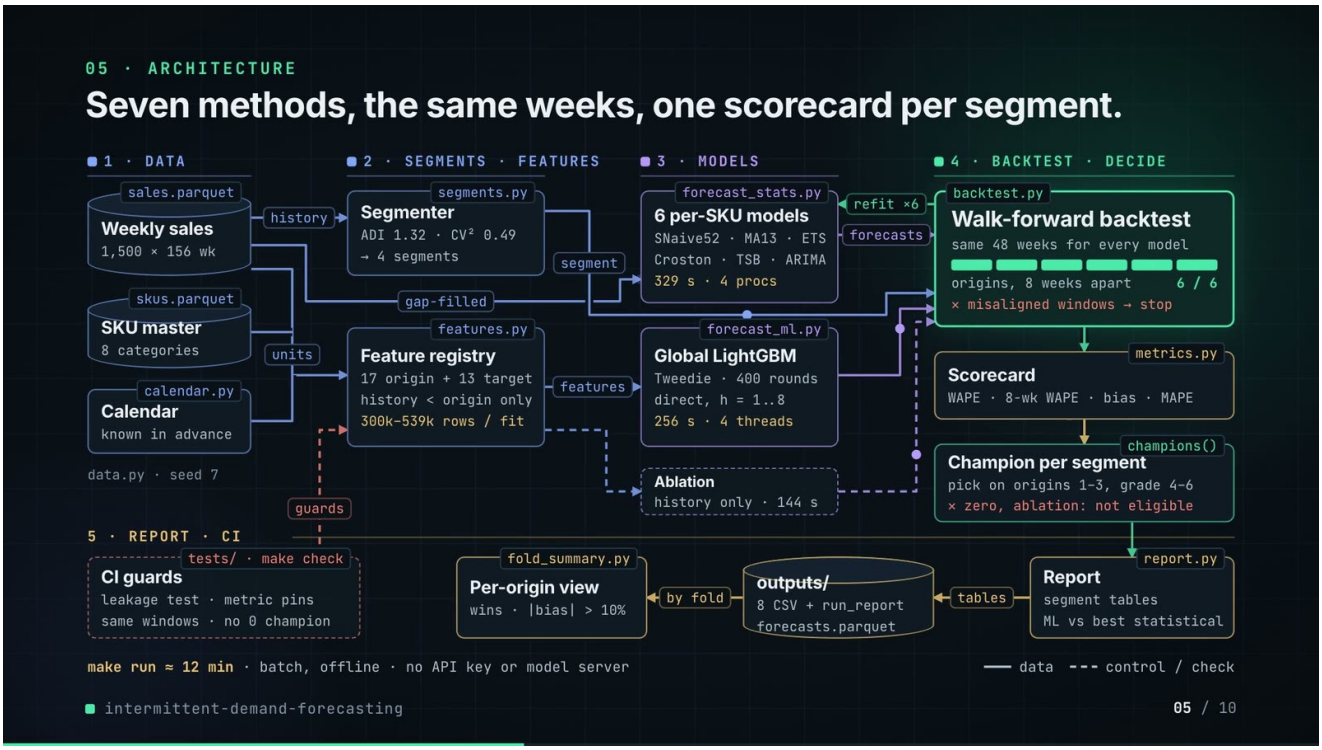
Illustrative, not data:

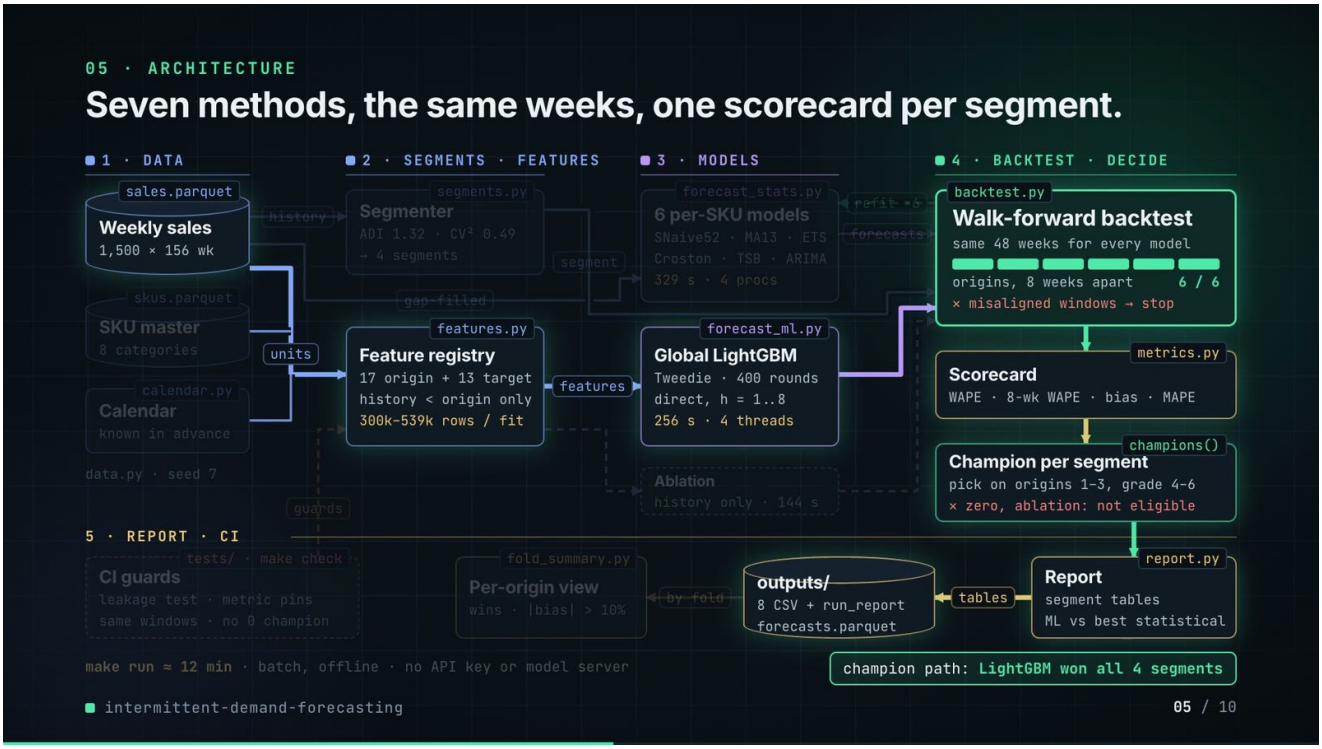
- The SOLD / ZERO / STEADY grid is the hand-worked case pinned in tests/test_metrics.py (F-12), not a real SKU.

To change this slide: Markup under #ex-b and the grid script in compositions/explainer.html; the footnote figures are F-08 and F-09.

Slide 05: Architecture

26.4-42.4 s · composition `compositions/s03-architecture.html`





On-screen text (verbatim)

05 · ARCHITECTURE

Seven methods, the same weeks, one scorecard per segment.

1 · DATA: sales.parquet Weekly sales (1,500 × 156 wk) · skus.parquet SKU master (8 categories) · calendar.py Calendar (known in advance) · data.py · seed 7

2 · SEGMENTS · FEATURES: segments.py Segmenter (ADI 1.32 · CV² 0.49 → 4 segments) · features.py Feature registry (17 origin + 13 target · history < origin only · 300k–539k rows / fit)

3 · MODELS: forecast_stats.py 6 per-SKU models (SNaive52 · MA13 · ETS · Croston · TSB · ARIMA · 329 s · 4 procs) · forecast_ml.py Global LightGBM (Tweedie · 400 rounds · direct, h = 1..8 · 256 s · 4 threads) · Ablation (history only · 144 s)

4 · BACKTEST · DECIDE: backtest.py Walk-forward backtest (same 48 weeks for every model · origins, 8 weeks apart 6 / 6 · ✕ misaligned windows → stop) → metrics.py Scorecard (WAPE · 8-wk WAPE · bias · MAPE) → champions() Champion per segment (pick on origins 1–3, grade 4–6 · ✕ zero, ablation: not eligible)

5 · REPORT · CI: report.py Report (segment tables · ML vs best statistical) → outputs/ (8 CSV + run_report · forecasts.parquet) → fold_summary.py Per-origin view (wins · |bias| > 10%) · tests/ · make check CI guards (leakage test · metric pins · same windows · no 0 champion)

Edge labels: history · units · gap-filled · features · segment · forecasts · refit ×6 · tables · by fold · guards

make run ≈ 12 min · batch, offline · no API key or model server · legend: data / control-check

champion path: LightGBM won all 4 segments

Figures on this slide

ID	Shown	Source · key	Calculation	Check
F-03	1,500 SKUs	outputs/sku_segments.csv · row count	-	✓
F-04	156 weeks / 1,500 × 156 wk	outputs/run_log.txt (printed by cli.py from data/sales.parquet) · '1500 SKUs, 156 weeks'	-	✓
F-07	7 methods, same windows / Seven methods	outputs/forecasts.parquet (model column) · models minus ZeroForecast and LightGBM-history-only	9 - 2 = 7	✓
F-11	REPLAY THE PAST · step forward, repeat · refit ×6	src/demand_segments/backtest.py:29-31; forecast_stats.py:49-57 · origins(); cross_validation(step_size=HORIZON)	-	✓
F-14	sales.parquet, skus.parquet, calendar.py, data.py, segments.py, features.py, forecast_stats.py, forecast_ml.py, backtest.py, metrics.py, champions(), report.py, fold_summary.py, tests/ · make check	src/demand_segments/, scripts/, tests/, Makefile · files + defs	-	✓
F-16	8 categories	src/demand_segments/data.py:20-29 (CATEGORIES) · dict entries	-	✓
F-17	data.py · seed 7	src/demand_segments/data.py:39 · seed: int = 7	-	✓
F-18	ADI 1.32 · CV² 0.49	src/demand_segments/segments.py:14-15 · ADI_CUT, CV2_CUT	-	✓
F-19	17 origin + 13 target	src/demand_segments/features.py (ORIGIN_FEATURES, TARGET_FEATURES) via demo assets/forecast_sheet.json arch · len(ORIGIN_FEATURES), len(TARGET_FEATURES)	-	✓
F-20	300k–539k rows / fit	features.training_frame(panel, origin) via demo assets/forecast_sheet.json arch · train_rows_first_origin = 299,648; train_rows_last_origin = 539,390	rounded to thousands	✓
F-21	SNaive52 · MA13 · ETS · Croston · TSB · ARIMA	src/demand_segments/forecast_stats.py:25-33 · models()	-	✓
F-22	329 s · 4 procs	outputs/run_report.md; Makefile:8 · 'statistical models: 329'; --jobs 4	-	✓
F-23	Tweedie · 400 rounds · direct, h = 1..8 · 256 s · 4 threads	src/demand_segments/forecast_ml.py:14, :25, :27; features.py:21; outputs/run_report.md · objective, num_threads, ROUNDS, HORIZON, 'LightGBM: 256'	-	✓
F-24	Ablation · history only · 144 s	outputs/run_report.md · 'LightGBM-history-only: 144'	-	✓
F-25	same 48 weeks for every model · origins, 8 weeks apart · 6 / 6	outputs/forecasts.parquet · distinct weeks; 6 folds	6 × 8 = 48	✓
F-26	✕ misaligned windows → stop	src/demand_segments/backtest.py:56-59 · RuntimeError	-	✓
F-27	WAPE · 8-wk WAPE · bias · MAPE	src/demand_segments/metrics.py · wape, window_wape, bias, mape	-	✓
F-28	pick on origins 1–3, grade 4–6	src/demand_segments/report.py:30; backtest.py:109 · n_select = n_windows // 2 = 3	-	✓
F-29	✕ zero, ablation: not eligible	src/demand_segments/backtest.py:108; tests/test_backtest.py:41 · isin([ZERO, ML_HISTORY_ONLY])	-	✓
F-30	8 CSV + run_report	src/demand_segments/report.py:41-48 · .to_csv calls	-	✓
F-31	wins · |bias| > 10%	scripts/fold_summary.py · origins_with_abs_bias_over_10pct	-	✓
F-32	CI guards · leakage test · metric pins · same windows · no 0 champion	tests/*.py; .github/workflows/ci.yml; Makefile check · test names	-	✓
F-33	make run ≈ 12 min / ~12 min	outputs/run_report.md 'Run time (seconds)' · 329 + 256 + 144 = 729 s	729 / 60 = 12.2 min	✓
F-34	batch, offline · no API key	README 'Running it' · 'Nothing needs an API key	-	✓

ID	Shown	Source · key	Calculation	Check
	or model server	or a model server'		
F-35	champion path: LightGBM won all 4 segments	`outputs/champions.csv` · champion	count(champion == LightGBM)	✓
F-36	refit ×6	`src/demand_segments/backtest.py:68-71 (fit per origin); forecast_stats.py:49-57 (cross_validation n_windows)` · for fold, origin in enumerate(starts)	-	✓

Assumptions: A-01, A-02, A-03, A-04, A-05, A-06, A-07, A-08, A-09

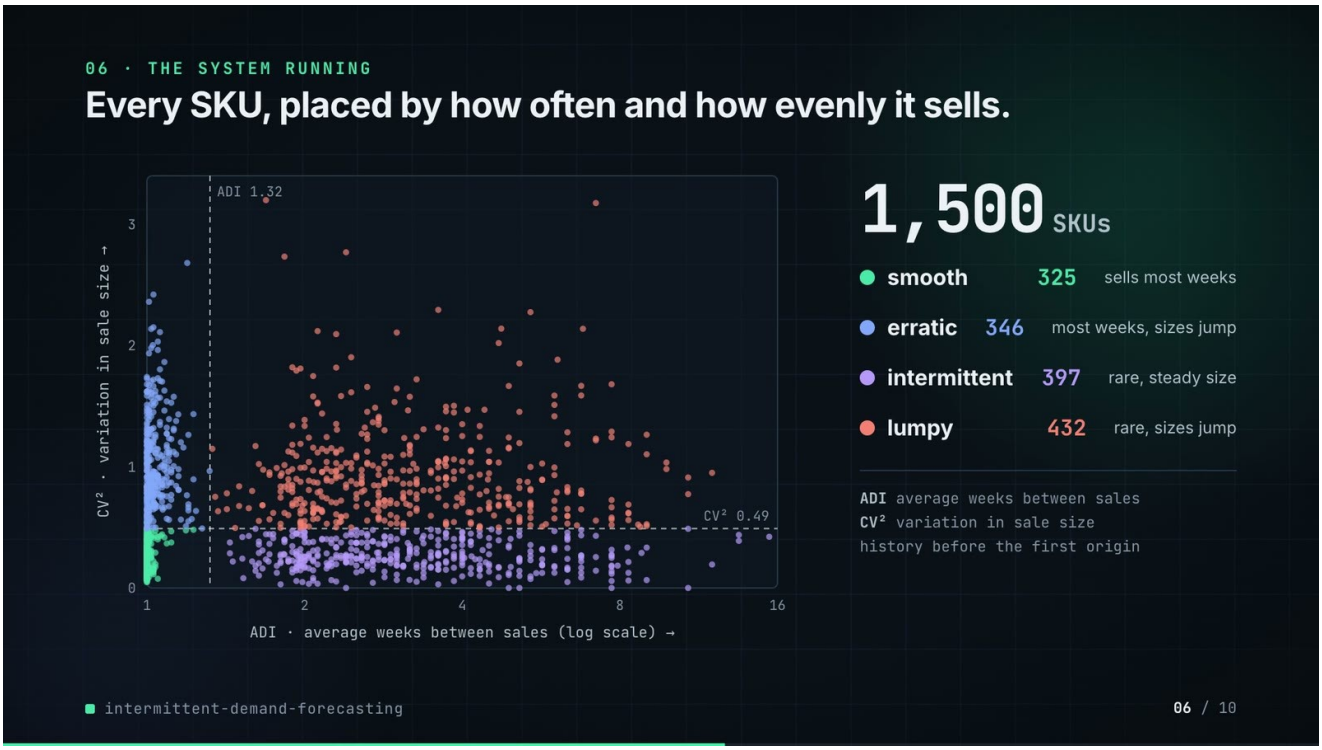
Illustrative, not data:

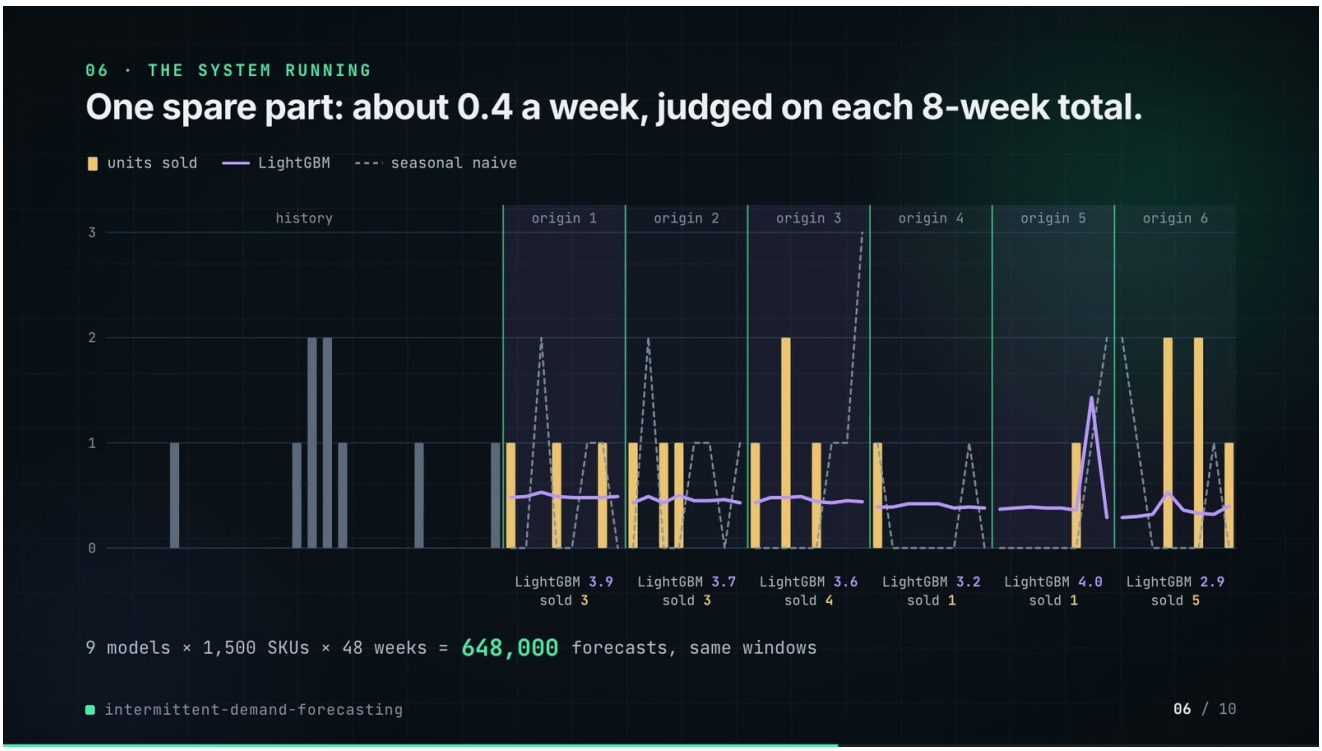
- Packets moving along the wires show execution order, not data volumes.
- The origin bar fills 1..6 in animation time, not real time.
- The 'champion path' highlight marks the LightGBM route through the pipeline; the other boxes are dimmed, not removed.

To change this slide: Node boxes (#n-*) and the W wires array (points, labels, execution order, crit flag) in compositions/s03-architecture.html.

Slide 06: The system running

42.0-56.6 s · composition `compositions/s04-routes.html (RUN block injected by tools/inject_sheet.py)`





On-screen text (verbatim)

06 · THE SYSTEM RUNNING

Every SKU, placed by how often and how evenly it sells.

ADI / CV² scatter with cut lines ADI 1.32, CV² 0.49 · 1,500 SKUs · smooth 325 sells most weeks · erratic 346 most weeks, sizes jump · intermittent 397 rare, steady size · lumpy 432 rare, sizes jump

ADI average weeks between sales · CV² variation in sale size · history before the first origin

One spare part: about 0.4 a week, judged on each 8-week total.

units sold · LightGBM · seasonal naive; history + origin 1..6; LightGBM 3.9 / sold 3 · 3.7 / 3 · 3.6 / 4 · 3.2 / 1 · 4.0 / 1 · 2.9 / 5

9 models × 1,500 SKUs × 48 weeks = 648,000 forecasts, same windows

Figures on this slide

ID	Shown	Source · key	Calculation	Check
F-03	1,500 SKUs	outputs/sku_segments.csv · row count	-	✓
F-18	ADI 1.32 · CV² 0.49	src/demand_segments/segments.py:14-15 · ADI_CUT, CV2_CUT	-	✓
F-25	same 48 weeks for every model · origins, 8 weeks apart · 6 / 6	outputs/forecasts.parquet · distinct weeks; 6 folds	6 × 8 = 48	✓
F-37	325 / 346 / 397 / 432	outputs/segment_counts.csv · skus	-	✓
F-38	1,500 dots on the ADI / CV² plane	outputs/sku_segments.csv; RUN block in compositions/s04-routes.html · adi, cv2, segment (rounded to 3 decimals)	-	✓
F-39	3.9/3 · 3.7/3 · 3.6/4 · 3.2/1 · 4.0/1 · 2.9/5	outputs/forecasts.parquet · sku_id = SKU00191, model = LightGBM, sum(yhat), sum(y) per fold	sums per fold, forecast rounded to 0.1	✓
F-40	about 0.4 a week	outputs/forecasts.parquet · SKU00191, LightGBM, mean(yhat) over 48 weeks = 0.44	-	✓
F-41	9 models × 1,500 SKUs × 48 weeks = 648,000 forecasts	outputs/forecasts.parquet · row count	9 × 1,500 × 48	✓
F-42	One spare part	outputs/sku_segments.csv; data/skus.parquet · SKU00191	-	✓

Assumptions: A-02, A-03, A-08

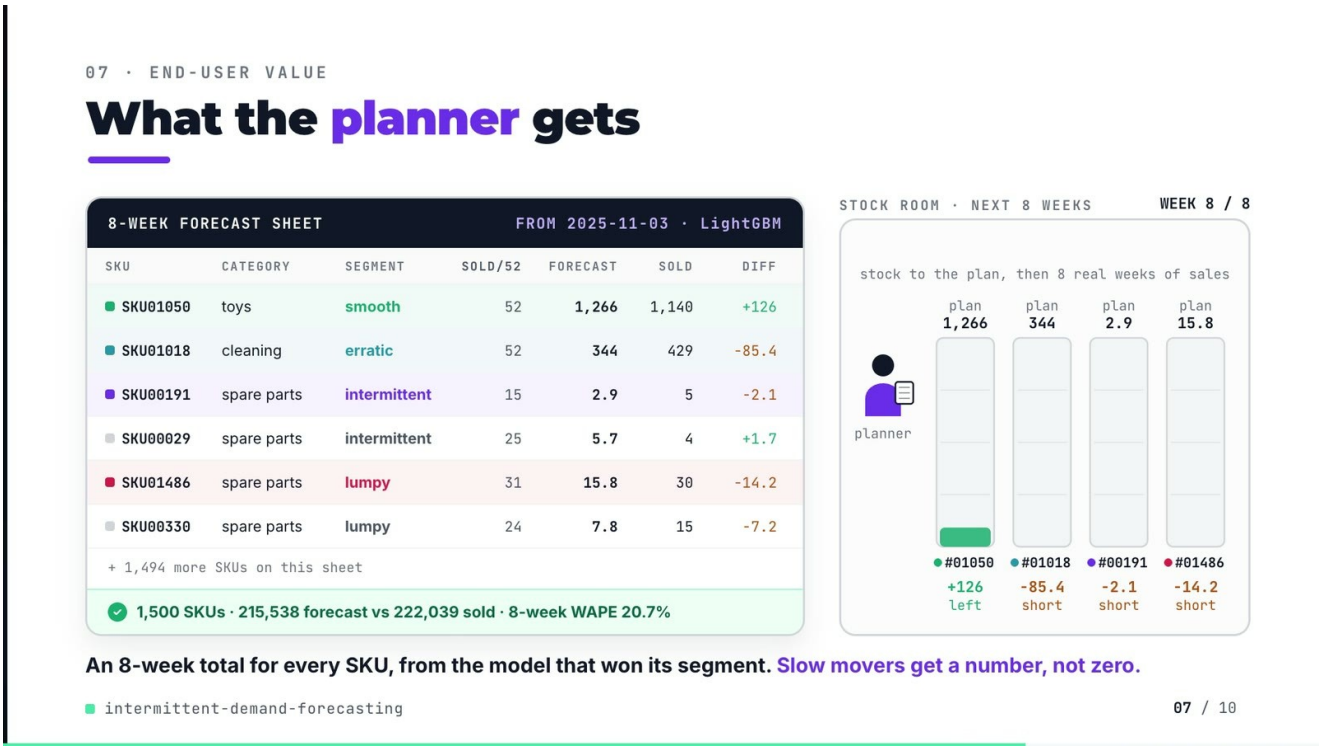
Illustrative, not data:

- The scatter is redrawn in SVG from outputs/sku_segments.csv (the repo saves no plot); CV² axis runs to 3.4 (max in data 3.20).
- The plain-words segment descriptions are the Syntetos-Boylan quadrant meanings.

To change this slide: Re-run tools/export_forecast_sheet.py + tools/inject_sheet.py --block RUN; layout in compositions/s04-routes.html.

Slide 07: End-user value (white)

56.3-69.3 s · composition `compositions/value.html` (SHEET block injected by `tools/inject_sheet.py`)



On-screen text (verbatim)

07 · END-USER VALUE

What the planner gets

8-WEEK FORECAST SHEET · FROM 2025-11-03 · LightGBM — columns SKU, CATEGORY, SEGMENT, SOLD/52, FORECAST, SOLD, DIFF

SKU01050 toys smooth 52 1,266 1,140 +126

SKU01018 cleaning erratic 52 344 429 −85.4

SKU00191 spare parts intermittent 15 2.9 5 −2.1

SKU00029 spare parts intermittent 25 5.7 4 +1.7

SKU01486 spare parts lumpy 31 15.8 30 −14.2

SKU00330 spare parts lumpy 24 7.8 15 −7.2

+ 1,494 more SKUs on this sheet · ✓ 1,500 SKUs · 215,538 forecast vs 222,039 sold · 8-week WAPE 20.7%

STOCK ROOM · NEXT 8 WEEKS · stock to the plan, then 8 real weeks of sales · planner · plan 1,266 / 344 / 2.9 / 15.8 · WEEK 8 / 8 · +126 left · −85.4 short · −2.1 short · −14.2 short

An 8-week total for every SKU, from the model that won its segment. Slow movers get a number, not zero.

Figures on this slide

ID	Shown	Source · key	Calculation	Check
F-03	1,500 SKUs	`outputs/sku_segments.csv` · row count	-	✓
F-43	SKU01050, SKU01018, SKU00191, SKU00029, SKU01486, SKU00330	`outputs/forecasts.parquet` (fold 6, LightGBM); `outputs/sku_segments.csv`; `data/skus.parquet`; `data/sales.parquet`; SHEET block in `compositions/value.html` · per-SKU sums	recomputed from the repo files, compared to the injected SHEET	✓
F-44	+ 1,494 more SKUs on this sheet	F-03 · 1,500 - 6 rows	1,500 - 6 = 1,494	✓
F-45	215,538 forecast vs 222,039 sold	`outputs/forecasts.parquet` · fold 6, LightGBM: sum(yhat), sum(y)	-	✓
F-46	8-week WAPE 20.7%	`outputs/forecasts.parquet` · fold 6, LightGBM	sum |8-wk forecast - 8-wk sold| / sum 8-wk sold over SKUs (metrics.window_wape)	✓
F-47	FROM 2025-11-03 · LightGBM	`outputs/forecasts.parquet` · first week of fold 6	-	✓

Assumptions: A-06, A-10, A-11

Illustrative, not data:

- Planner figure, boxes and bays are drawings; bay heights are scaled per bay.
- The bays drain by each week's real units sold; this is forecast minus actual, not a service-level or inventory simulation (A-10).

- 'Slow movers get a number, not zero' describes the LightGBM forecasts shown; it is framing, not a measured result.

To change this slide: Re-run tools/export_forecast_sheet.py, then tools/inject_sheet.py --block SHEET; layout and scene in compositions/value.html.

Slide 08: Results

69.0-77.8 s · composition `compositions/s05-results.html (stats array)`



On-screen text (verbatim)

08 · RESULTS · 1,500 SKUS · 6 ORIGINS

LightGBM wins all four segments on the 8-week total.

LIGHTGBM 21.9% 8-week WAPE, all SKUs · seasonal naive 29.1%

CHAMPION · SMOOTH −25.8% weekly WAPE vs seasonal naive · 37.6% → 27.9%

CHAMPION · ERRATIC −30.7% weekly WAPE vs seasonal naive · 74.8% → 51.9%

BACKTEST 24/24 segment-origins won on 8-week WAPE · 4 segments × 6 origins

champions picked on origins 1–3 and graded on 4–6 · cards 2–3 are origins 4–6

Figures on this slide

ID	Shown	Source · key	Calculation	Check
F-03	1,500 SKUs	`outputs/sku_segments.csv` · row count	-	✓
F-28	pick on origins 1–3, grade 4–6	`src/demand_segments/report.py:30`; `backtest.py:109` · n_select = n_windows // 2 = 3	-	✓
F-50	21.9%	`outputs/scores_by_segment.csv` · segment=all, model=LightGBM, wape_8wk	x 100	✓
F-51	seasonal naive 29.1%	`outputs/scores_by_segment.csv` · segment=all, model=SeasonalNaive, wape_8wk	x 100	✓
F-52	-25.8%	`outputs/champions.csv` · segment=smooth, wape_reduction_vs_seasonal_naive	-(x 100)	✓
F-53	37.6% → 27.9%	`outputs/champions.csv` · seasonal_naive_wape, champion_wape	-	✓
F-54	-30.7%	`outputs/champions.csv` · segment=erratic, wape_reduction_vs_seasonal_naive	-(x 100)	✓
F-55	74.8% → 51.9%	`outputs/champions.csv` · seasonal_naive_wape, champion_wape	-	✓
F-56	24/24 · 4 segments × 6 origins	`outputs/fold_summary.csv` · sum(origins_won_on_wape_8wk), sum(origins)	-	✓
F-57	LightGBM wins all four segments on the 8-week total.	`outputs/scores_by_segment.csv` · argmin(wape_8wk) excluding ZeroForecast and the ablation	-	✓

Assumptions: A-04, A-05, A-06, A-11

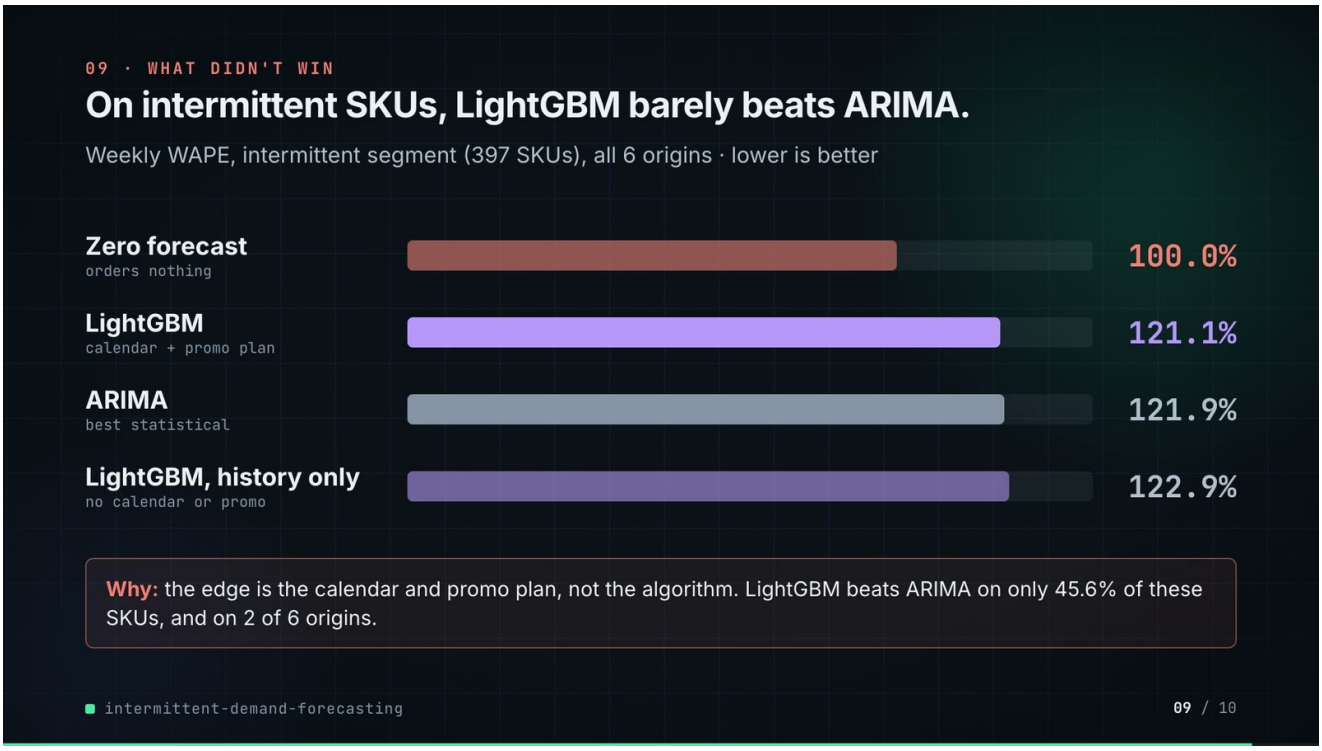
Illustrative, not data:

- Mini bars under each number are drawn to the same ratio (after / before; 24/24 = full).

To change this slide: The stats array and card text in compositions/s05-results.html.

Slide 09: What didn't win

77.4-84.4 s · composition `compositions/s06-twist.html`



On-screen text (verbatim)

09 · WHAT DIDN'T WIN

On intermittent SKUs, LightGBM barely beats ARIMA.

Weekly WAPE, intermittent segment (397 SKUs), all 6 origins · lower is better

Zero forecast (orders nothing) 100.0% · LightGBM (calendar + promo plan) 121.1% · ARIMA (best statistical) 121.9% · LightGBM, history only (no calendar or promo) 122.9%

Why: the edge is the calendar and promo plan, not the algorithm. LightGBM beats ARIMA on only 45.6% of these SKUs, and on 2 of 6 origins.

Figures on this slide

ID	Shown	Source · key	Calculation	Check
F-58	100.0%	`outputs/scores_by_segment.csv` · segment=intermittent, model=ZeroForecast, wape	x 100	✓
F-59	121.1%	`outputs/scores_by_segment.csv` · segment=intermittent, model=LightGBM, wape	x 100	✓
F-60	121.9%	`outputs/scores_by_segment.csv`; `outputs/lightgbm_vs_best_statistical.csv` · segment=intermittent, model=AutoARIMA, wape; best_statistical = AutoARIMA	x 100	✓
F-61	122.9%	`outputs/scores_by_segment.csv` · segment=intermittent, model=LightGBM-history-only, wape	x 100	✓
F-62	397 SKUs	`outputs/segment_counts.csv` · segment=intermittent	-	✓
F-63	45.6%	`outputs/lightgbm_vs_best_statistical.csv` · segment=intermittent, lightgbm_sku_win_share	x 100	✓
F-64	2 of 6 origins	`outputs/fold_summary.csv` · segment=intermittent, origins_won_on_wape / origins	-	✓
F-65	the edge is the calendar and promo plan, not the algorithm	`outputs/scores_by_segment.csv` · intermittent: LightGBM-history-only 122.9% > AutoARIMA 121.9% > LightGBM 121.1%	-	✓

Assumptions: A-05, A-07, A-11

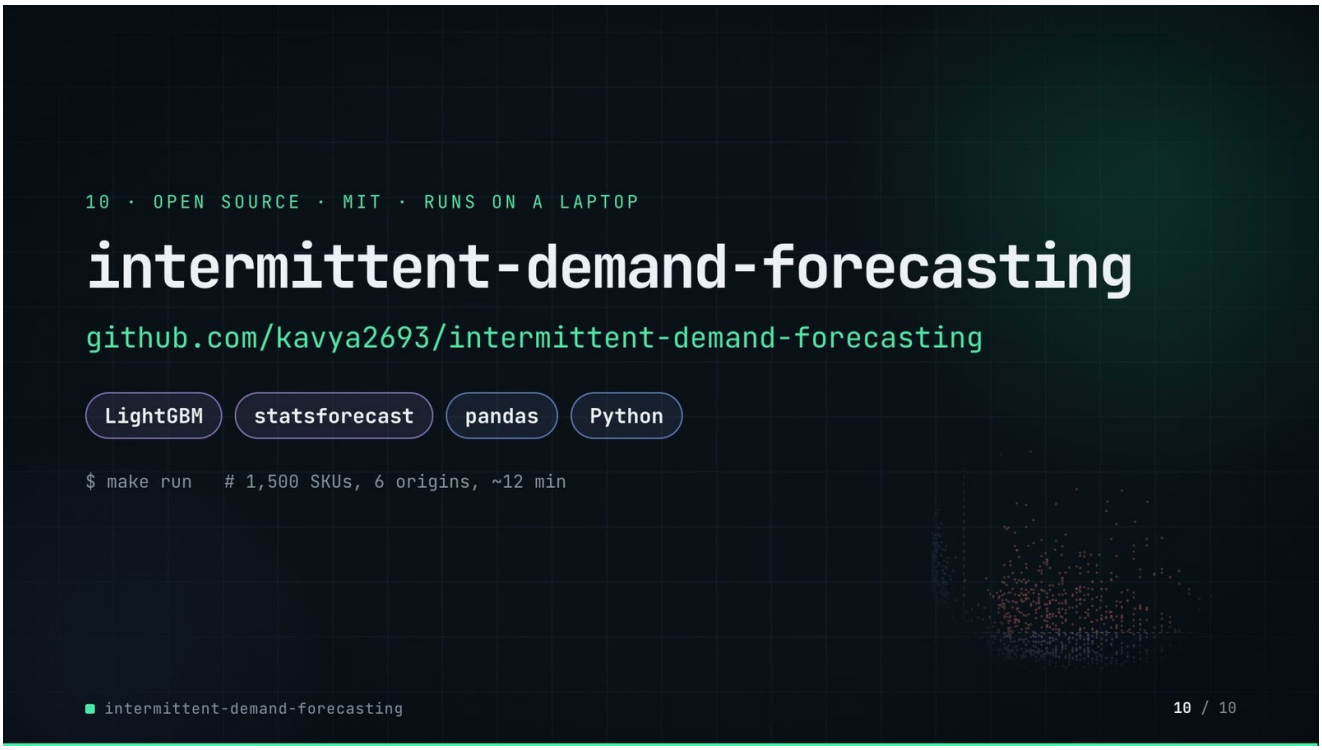
Illustrative, not data:

- Bar lengths: 1,000 px track = 140% weekly WAPE.

To change this slide: Rows #twist-r1..r4 (bar widths are inline styles) in compositions/s06-twist.html.

Slide 10: End card

84.0-88.4 s · composition `compositions/s07-end.html`



On-screen text (verbatim)

10 · OPEN SOURCE · MIT · RUNS ON A LAPTOP
intermittent-demand-forecasting
github.com/kavya2693/intermittent-demand-forecasting
LightGBM · statsforecast · pandas · Python
$ make run # 1,500 SKUs, 6 origins, ~12 min

Figures on this slide

ID	Shown	Source · key	Calculation	Check
F-33	make run ≈ 12 min / ~12 min	`outputs/run_report.md 'Run time (seconds)'` · 329 + 256 + 144 = 729 s	729 / 60 = 12.2 min	✓
F-66	MIT	`LICENSE line 1` · -	-	✓
F-67	$ make run # 1,500 SKUs, 6 origins, ~12 min	`Makefile:7-9; README 'Running it'` · run target	-	✓
F-68	LightGBM · statsforecast · pandas · Python	`pyproject.toml dependencies` · lightgbm, statsforecast, pandas	-	✓

Assumptions: A-09, A-11

Illustrative, not data:

- Faded scatter in the background is outputs/sku_segments.csv drawn by tools/make_end_scatter.py.

To change this slide: Text in compositions/s07-end.html.

5. Change-impact matrix

Input / assumption	Affected figures	Slides	Files to edit	Re-run
A-01: All data is synthetic, produced by a seeded generator sized to the production problem: 1,500 SKUs, 156 weekly periods, 8 categories, B2B and B2C channels, seed 7	all figures	01, 02, 03, 04, 05, 06, 07, 08, 09, 10	`src/demand_segments/data.py (GeneratorConfig, CATEGORIES, PATTERN_MIX)`	make data; make run; export_forecast_sheet.py + inject_sheet.py
A-02: Segments use the Syntetos-Boylan cut-offs ADI 1	F-05, F-06, F-08, F-09, F-18, F-37, F-38, F-62	02, 04, 05, 06, 09	`src/demand_segments/segments.py:14-15`	make run; export_forecast_sheet.py + inject_sheet.py
A-03: Walk-forward backtest: six origins, eight weeks apart, each forecasting the next eight weeks (48 test weeks), the same windows for every model	F-07, F-11, F-25, F-36, F-39, F-41, F-56, F-64	02, 03, 05, 06, 08, 09	`src/demand_segments/cli.py:18; src/demand_segments/features.py:21`	make run; export_forecast_sheet.py + inject_sheet.py
A-04: Each segment's champion is picked on eight-week WAPE over origins 1-3 and graded on origins 4-6; the zero forecast and the history-only ablation are never eligible	F-10, F-28, F-29, F-35, F-52, F-53, F-54, F-55	03, 05, 08	`src/demand_segments/backtest.py:100-130; src/demand_segments/report.py:30`	make run
A-05: ETS and ARIMA are fitted without a 52-week season (season_length=1); only seasonal naive and LightGBM carry the yearly pattern, so the statistical side is weaker on seasonal SKUs than a tuned seasonal model would be	F-21, F-53, F-55, F-58, F-59, F-60, F-61, F-63, F-64	05, 08, 09	`src/demand_segments/forecast_stats.py:22-33`	make run
A-06: LightGBM settings: Tweedie objective (variance power 1	F-23, F-50, F-52, F-54, F-56, F-57, F-59	05, 08, 09	`src/demand_segments/forecast_ml.py:11-27`	make run; export_forecast_sheet.py + inject_sheet.py
A-07: LightGBM may read the calendar, the promotion plan and the price list for the target week (11 'known-future' features), because a planner has them in advance	F-24, F-61, F-65	05, 09	`src/demand_segments/forecast_ml.py:30-42; src/demand_segments/features.py (TARGET_FEATURES)`	make run
A-08: Feed-gap weeks are filled with a trailing 8-week mean for the statistical models' inputs only; every score uses the raw actuals, and missing actual weeks are skipped, never scored as zero	F-41, F-45, F-46	06, 07	`src/demand_segments/forecast_stats.py:37-46`	make run
A-09: Run times are wall-clock seconds from one make run on the author's laptop with --jobs 4 (statistical models) and 4 LightGBM threads; they change with the machine and its load	F-22, F-23, F-24, F-33, F-67	05, 10	`Makefile:8`	make run
A-10: Slide 07 shows the LAST backtest origin (2025-11-03, weeks to 2025-12-22) as if it were the planner's sheet, with what actually sold next to it	F-43, F-44, F-45, F-46, F-47	07	`tools/export_forecast_sheet.py (pick rules)`	export_forecast_sheet.py + inject_sheet.py; hyperframes render
A-11: Video rounding: percentages to 1 decimal, units and counts to whole numbers with thousands separators, 8-week totals under 100 units to 1 decimal; '300k-539k', '≈ 12 min' and 'about 0	F-20, F-33, F-40, F-45, F-46, F-50, F-51, F-52, F-53, F-54, F-55, F-58, F-59, F-60, F-61, F-63, F-67	05, 06, 07, 08, 09, 10	the composition file named on each slide	hyperframes render
A-12: Slide 01 (industries, everyday problems) is general demand-planning knowledge used as context	F-01, F-02	01	`compositions/s01-context.html`	hyperframes render

6. Known differences and caveats

• All results are on generated data (A-01): they show the method behaving as expected on realistic patterns, not a measured business impact (README 'Known limitations').

• Forecasts are point forecasts only; slide 07's 'left / short' is forecast minus actual over 8 weeks, not a fill-rate or stock simulation (A-10).

• ETS and ARIMA run without a yearly season (A-05), which makes the statistical side weaker on seasonal SKUs than a tuned seasonal model.

• LightGBM under-forecasts lumpy SKUs (bias -8.9% over all origins, -15.7% at its worst origin); the video does not show this, the README does.

• Weekly WAPE above 100% on the intermittent and lumpy segments is expected: a zero forecast scores exactly 100%, which is why the scorecard uses the 8-week total.

• Slide 07 rows are median-error SKUs per segment chosen by the demo export script, not by the repo.

• Run times (A-09) are one run on the author's laptop.

• Slide 01 statements are general demand-planning knowledge (A-12).

Appendix: environment

Item	Value
Commit	`3616b6c2747b6be20b0c56b6345c4954391b22ac`
Python	3.11.7
lightgbm	4.7.0
statsforecast	2.0.1
pandas	3.0.6
numpy	2.4.6
Profile / seeds	data_seed=7, lightgbm_seed=7, skus=1500, weeks=156, origins=6, horizon_weeks=8
Run time	outputs/run_report.md: statistical models 329 s + LightGBM 256 s + history-only ablation 144 s = 729 s (about 12 min) for make run (--jobs 4) on the author's laptop